Why AI for PM?

- PM work contains repetitive analysis and reporting tasks.
- AI can accelerate drafting, but project decisions remain human-owned.
- The pilot focuses on measurable value, not novelty.

Three supported use cases

- 1. Risk & Issue Analysis
- 2. Executive Status Reporting
- 3. Meeting -> Actions & Decisions

The harness

- Raw data -> Anonymization -> Task context -> AI -> Validation -> Human review -> Approved output

What can go wrong?

- Hallucinated owner
- Invented date
- Unsupported risk
- Confidentiality leak
- Wrong confidence / overstatement

Evidence grounding

- Every material claim should carry a source identifier.
- No evidence -> UNKNOWN or UNSUPPORTED_CLAIM.

Human-in-the-loop

- AI creates a draft.
- PM reviews factual correctness, business interpretation, risk severity and publication decision.

How we measure value

- Accuracy / F1
- Unsupported claim rate
- Confidentiality leakage
- Median time saved
- Human correction effort

Choosing the model

- Compare models on the same cases and the same schema.
- Choose per task when evidence supports task-specific routing.

Safe usage rules

- Minimize data.
- Never paste credentials.
- Do not let AI make final project decisions.
- Review high-impact outputs.

Takeaway

- AI augments PM execution under controlled governance.